

Monte Carlo Search

Tristan Cazenave

LAMSADE CNRS

Université Paris-Dauphine

PSL

Tristan.Cazenave@dauphine.fr

Outline

- Monte Carlo Tree Search
- Nested Monte Carlo Search
- Nested Rollout Policy Adaptation
- Playout Policy Adaptation
- Imperfect Information Games
- Zero Learning (Deep RL)

INTELLIGENCE ARTIFICIELLE

Une approche ludique

Tristan Cazenave



Monte Carlo Tree Search

Monte Carlo Go

- 1993 : first Monte Carlo Go program
 - Gobble, Bernd Bruegmann.
 - How nature would play Go ?
 - Simulated annealing on two lists of moves.
 - Statistics on moves.
 - Only one rule : do not fill eyes.
 - Result = average program for 9x9 Go.
 - Advantage : much more simple than alternative approaches.

Monte Carlo Go

- 1998 : first master course on Monte Carlo Go.
- 2000 : sampling based algorithm instead of simulated annealing.
- 2001 : Computer Go an AI Oriented Survey.
- 2002 : Bernard Helmstetter.
- 2003 : Bernard Helmstetter, Bruno Bouzy, Developments on Monte Carlo Go.

Monte Carlo Phantom Go

- Phantom Go is Go when you cannot see the opponent's moves.
- A referee tells you illegal moves.
- 2005 : Monte Carlo Phantom Go program.
- Many Gold medals at computer Olympiad since then using flat Monte Carlo.
- 2011 : Exhibition against human players at European Go Congress.

UCT

- UCT : Exploration/Exploitation dilemma for trees [Kocsis and Szepesvari 2006].
- Play random random games (playouts).
- Exploitation : choose the move that maximizes the mean of the playouts starting with the move.
- Exploration : add a regret term (UCB).

UCT


- UCT : exploration/exploitation dilemma.
- Play the move that maximizes

$$\frac{w_i}{n_i} + c \sqrt{\frac{\ln t}{n_i}}$$

In which

- w_i = number of wins after the i -th move
- n_i = number of simulations after the i -th move
- c = exploration parameter (theoretically equal to $\sqrt{2}$)
- t = total number of simulations for the parent node

UCT




1) descent of the tree

2) playout


3) update the tree

UCT


Selection




Expansion



Sampling



Backpropagation



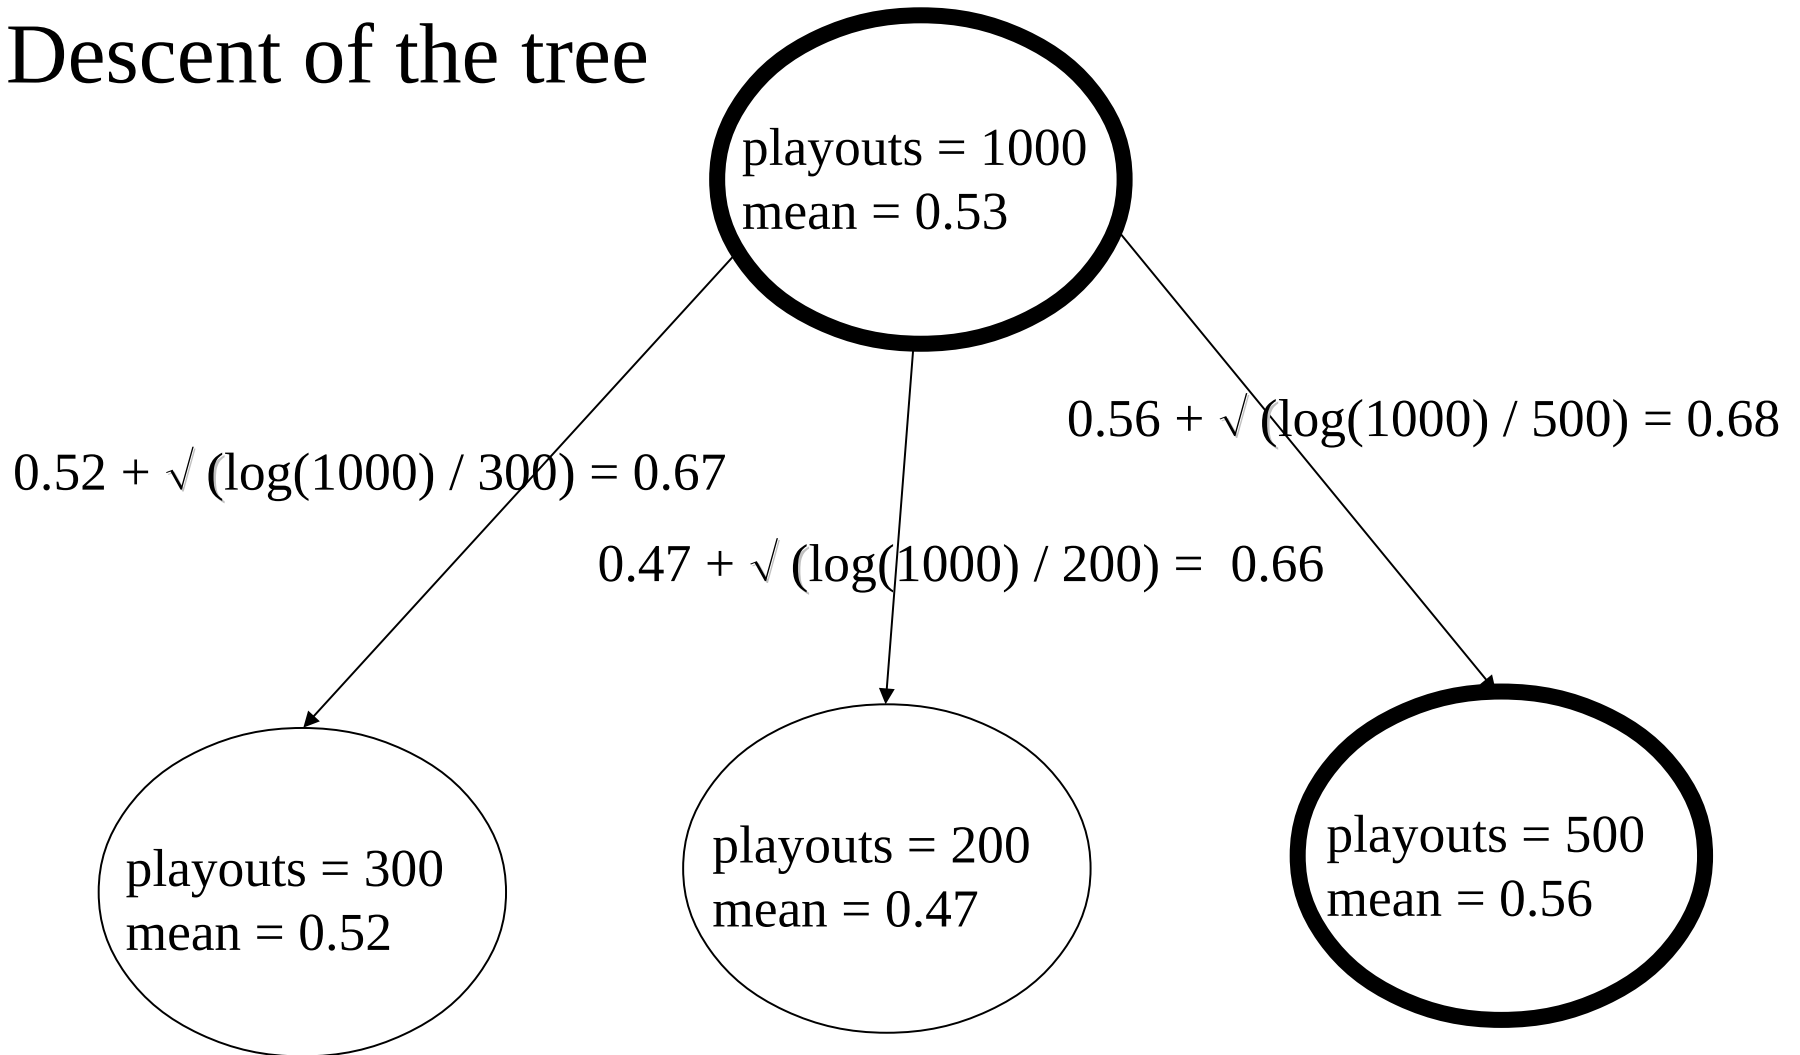
Tree Policy

Default Policy



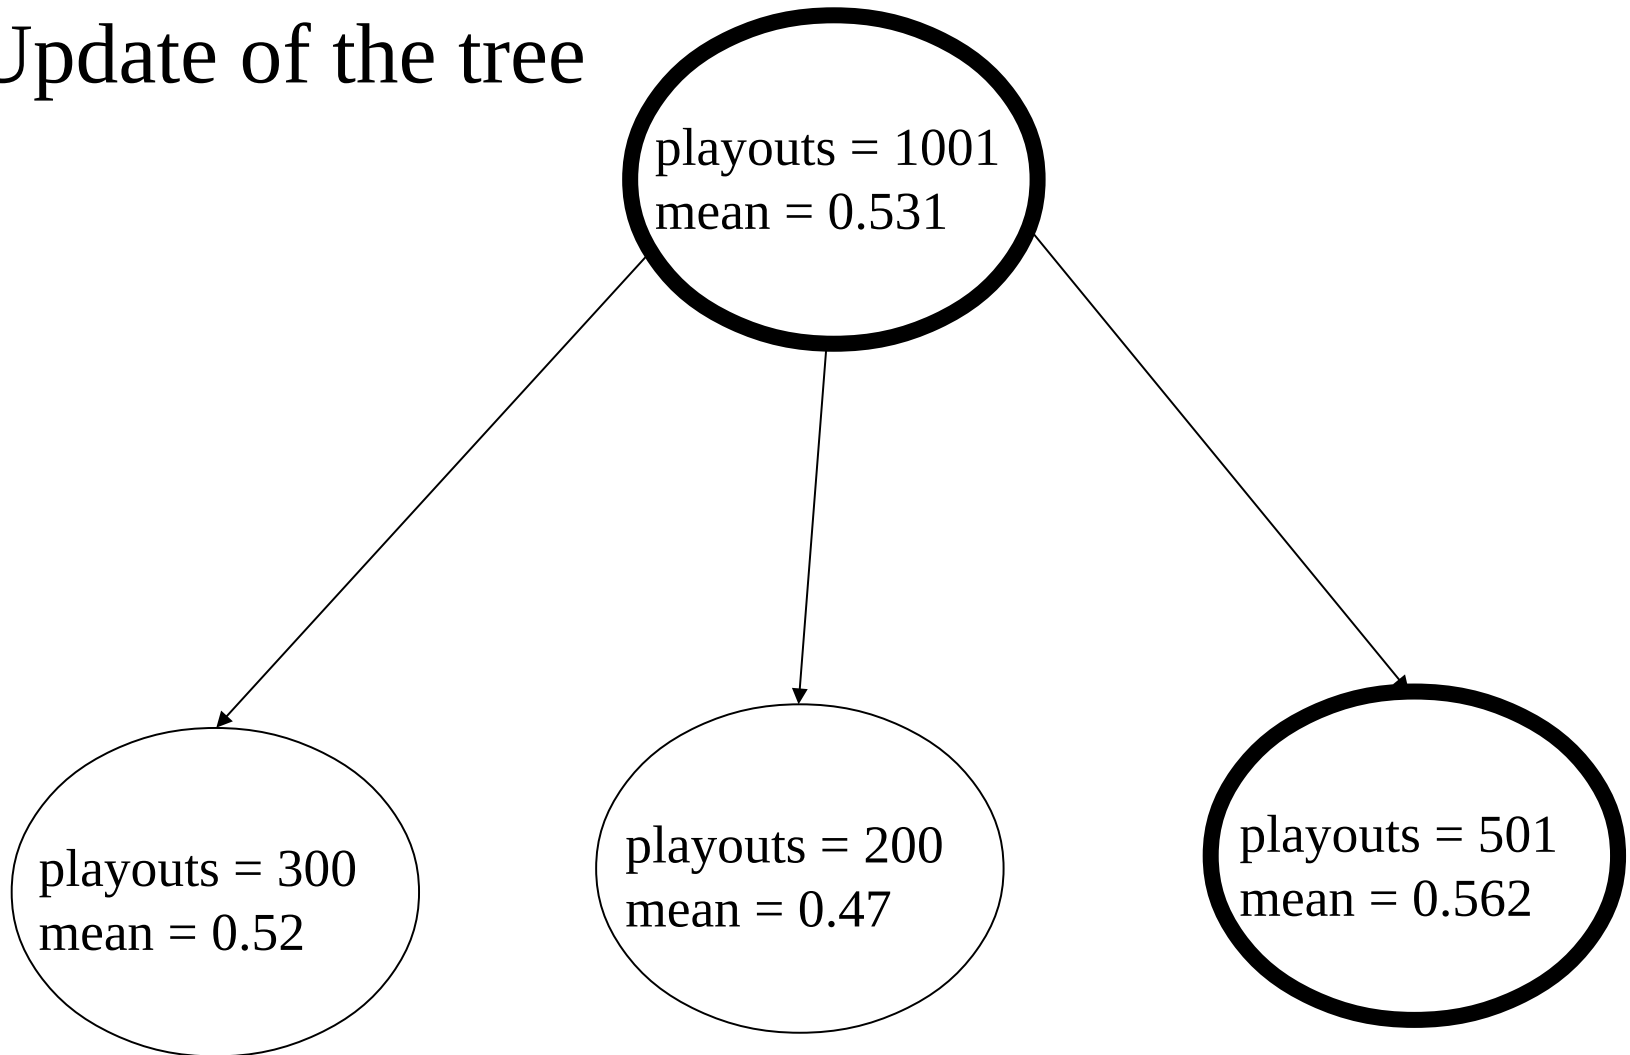
UCT

Descent of the tree



UCT

Update of the tree



UCT

```
function MCTSSEARCH( $s_0$ )  
  create root node  $v_0$  with state  $s_0$   
  while within computational budget do  
     $v_l \leftarrow$  TREEPOLICY( $v_0$ )  
     $\Delta \leftarrow$  DEFAULTPOLICY( $s(v_l)$ )  
    BACKUP( $v_l, \Delta$ )  
  return  $a(\text{BESTCHILD}(v_0, 0))$ 
```

UCT

```
function TREEPOLICY( $v$ )  
  while  $v$  is nonterminal do  
    if  $v$  not fully expanded then  
      return EXPAND( $v$ )  
    else  
       $v \leftarrow$  BESTCHILD( $v, Cp$ )  
  return  $v$ 
```

UCT

```
function EXPAND( $v$ )  
  choose  $a \in$  untried actions from  $A(s(v))$   
  add a new child  $v'$  to  $v$   
    with  $s(v') = f(s(v), a)$   
    and  $a(v') = a$   
return  $v'$ 
```


UCT

```
function BESTCHILD( $v, c$ )  
  return  $\arg \max_{v' \in \text{children of } v} \frac{Q(v')}{N(v')} + c \sqrt{\frac{2 \ln N(v)}{N(v)}}$ 
```

UCT

```
function DEFAULTPOLICY( $s$ )  
  while  $s$  is non-terminal do  
    choose  $a \in A(s)$  uniformly at random  
     $s \leftarrow f(s, a)$   
  return reward for state  $s$ 
```

UCT

```
function BACKUP( $v, \Delta$ )  
  while  $v$  is not null do  
     $N(v) \leftarrow N(v) + 1$   
     $Q(v) \leftarrow Q(v) + \Delta(v, p)$   
     $v \leftarrow$  parent of  $v$ 
```

AMAF

- All Moves As First (AMAF).
- AMAF calculates for each possible move of a state the average of the playouts that contain this move.

RAVE

- A big improvement for Go, Hex and other games is Rapid Action Value Estimation (RAVE) [Gelly and Silver 2007].
- RAVE combines the mean of the playouts that start with the move and the mean of the playouts that contain the move.

RAVE

- Parameter β_m for move m is :

$$\beta_m \leftarrow \text{pAMAF}_m /$$
$$(\text{pAMAF}_m + p_m + \text{bias} \times \text{pAMAF}_m \times p_m)$$

- β_m starts at 1 when no playouts and decreases as more playouts are played.
- Selection of moves in the tree :


$$\text{argmax}_m ((1.0 - \beta_m) \times \text{mean}_m + \beta_m \times \text{AMAF}_m)$$

GRAVE

- Generalized Rapid Action Value Estimation (GRAVE) is a simple modification of RAVE.
- It consists in using the first ancestor node with more than n playouts to compute the RAVE values.
- It is a big improvement over RAVE for Go, Atarigo, Knightthrough and Domineering [Cazenave 2015].

Atarigo

- Atarigo is a simplification of the game of Go.
- The winner of a game is the first player to capture a string of stones.




White to play and win.

Atarigo

- Atarigo 8x8 with 10,000 playouts :
RAVE wins 94.2 % against UCT.
GRAVE wins 88.4 % against RAVE.
- Atarigo 19x19 with 1,000 playouts :
RAVE wins 72.4 % against UCT.
GRAVE wins 78.2 % against RAVE.

Knighttthrough



Knightthrough


- RAVE wins 69.4 % against UCT with 1,000 playouts.
- GRAVE wins 67.8 % against RAVE with 1,000 playouts.
- RAVE wins 56.2 % against UCT with 10,000 playouts.
- GRAVE wins 67.2 % against RAVE with 10,000 playouts.

Domineering

- Domineering is a two player combinatorial game usually played on an 8x8 board.
- It consists in playing 2x1 dominoes on the board.
- The first player put the dominoes vertically and the second player put them horizontally.
- If a player cannot play anymore, he loses the game.

Domineering

- The last to play has won.



Domineering

- Domineering 8x8 with 10,000 playouts :
RAVE wins 72.6 % against UCT.
GRAVE wins 62.4 % against RAVE.
- Domineering 19x19 with 1,000 playouts :
RAVE wins 63.8 % against UCT.
GRAVE wins 56.4 % against RAVE.

Go

- Go is an ancient oriental game of strategy that originated in China thousands of years ago.
- It is usually played on a 19x19 grid.
- AlphaGo is the current best (computer) Go player.
- MCTS is the best algorithm for the game of Go.
- The RAVE algorithm was originally designed for computer Go.

Go 9x9

- Go 9x9 with 1,000 playouts :
RAVE wins 89.6 % against UCT.
GRAVE wins 66.0 % against RAVE.
- Go 9x9 with 10,000 playouts :
RAVE wins 73.2 % against UCT.
GRAVE wins 54.4 % against RAVE.

Go 19x19

- Go 19x19 with 1,000 playouts :
GRAVE wins 81.8 % against RAVE.
- Go 19x19 with 10,000 playouts :
GRAVE wins 62.4 % against RAVE.

Three Color Go

- Multicolor Go is Go with more than two players.
- Three Color Go is played with stones of three different colors.
- Chinese rules are used to score games.
- The winner is the player that has the greatest score at the end.
- The average winning rate is 33.33 %.

Three Color Go

- Three Color Go 9x9 with 1,000 playouts :
RAVE wins 70.83 % against UCT.
GRAVE wins 57.17 % against RAVE.
- Three Color Go 19x19 with 1,000 playouts :
RAVE wins 18.50 % against two GRAVE.

Parallelization of MCTS

- Root Parallelization.
- Tree Parallelization (virtual loss).
- Leaf Parallelization.

MCTS



- Great success for the game of Go since 2007.
- Much better than all previous approaches to computer Go.

AlphaGo

Lee Sedol is among the strongest and most famous 9p Go player :



AlphaGo has won 4-1 against Lee Sedol in March 2016
AlphaGo Master wins 3-0 against Ke Jie, 60-0 against pros.
AlphaGo Zero wins 89-11 against AlphaGo Master in 2017.

General Game Playing



- General Game Playing = play a new game just given the rules.
- Competition organized every year by Stanford.
- Ary world champion in 2009 and 2010.
- All world champions since 2007 use MCTS.




General Game Playing



- Eric Piette combined Stochastic Constraint Programming with Monte Carlo in WoodStock.
- World champion in 2016 (MAC-UCB-SYM).
- Detection of symmetries in the states.

Other two-player games


- Hex : 2009
- Amazons : 2009
- Lines of Action : 2009




MCTS Solver

- When a subtree has been completely explored the exact result is known.
- MCTS can solve games.
- Score Bounded MCTS is the extension of pruning to solving games with multiple outcomes.

Breakthrough



(a) Starting position on size 5×5 .



(b) Possible movements.

- Write the Board and Move classes for Breakthrough 5x5.
- Write the function for the possible moves.
- Write a program to play random games at Breakthrough 5x5.

Breakthrough

- The Move class contains the color, the starting and arriving locations of a pawn.

```
class Move(object):
```

```
    def __init__(self, color, x1, y1, x2, y2):
```

```
        self.color = color
```

```
        self.x1 = x1
```

```
        self.y1 = y1
```

```
        self.x2 = x2
```

```
        self.y2 = y2
```

Breakthrough

- The Board class initializes the board with two rows of Black and two rows of White pawns:

Dx = 5

Dy = 5

Empty = 0

White = 1

Black = 2

```
class Board(object):
```

```
    def __init__(self):
```

```
        self.h = 0
```

```
        self.turn = White
```

```
        self.board = np.zeros ((Dx, Dy))
```

```
        for i in range (0, 2):
```

```
            for j in range (0, Dy):
```

```
                self.board [i] [j] = White
```

```
        for i in range (Dx - 2, Dx):
```

```
            for j in range (0, Dy):
```

```
                self.board [i] [j] = Black
```

Breakthrough

- Test if a move is valid for a given board:

```
def valid (self, board):
    if self.x2 >= Dx or self.y2 >= Dy or self.x2 < 0 or self.y2 < 0:
        return False
    if self.color == White:
        if self.x2 != self.x1 + 1:
            return False
        if board.board [self.x2] [self.y2] == Black:
            if self.y2 == self.y1 + 1 or self.y2 == self.y1 - 1:
                return True
            return False
        elif board.board [self.x2] [self.y2] == Empty:
            if self.y2 == self.y1 + 1 or self.y2 == self.y1 - 1 or self.y2 == self.y1:
                return True
            return False
    elif self.color == Black:
        ...
```

Breakthrough

- Generate the legal moves:

```
def legalMoves(self, color):
```

```
moves = []
```

```
for i in range (0, Dx):
```

```
for j in range (0, Dy):
```

```
if self.board [i] [j] == color:
```

for k in $[-1, 0, 1]$:

```
for l in [-1, 0, 1]:
```

m = Move (color, i, j, i + k, j + l)

```
if m.valid (self):
```

moves.append (m)

return moves

Flat Monte Carlo

- Write a won function to detect when a game is won.
- Write a playout function that plays a random game from the current state and returns the result of the random game (1.0 if White wins, 0.0 else).

Flat Monte Carlo

- Keep statistics for all the moves of the initial state.
- For each move of the initial state, keep the number of playouts starting with the move and the number of playouts starting with the move that have been won.
- Play the move with the greatest mean when all the playouts are finished.

UCB

Choose the first move at the root according to UCB before each playout:

$$\frac{w_i}{n_i} + c \sqrt{\frac{\ln t}{n_i}}$$

In which

- w_i = number of wins after the i -th move
- n_i = number of simulations after the i -th move
- c = exploration parameter (theoretically equal to $\sqrt{2}$)
- t = total number of simulations for the parent node

UCB vs Flat

- Make UCB with 1 000 playouts play 100 games against Flat with 1 000 playouts.
- Tune the UCB constant (hint 0.4).